# CS 179: GPU Computing

**Recitation 2**: Synchronization, Shared memory, Matrix Transpose

# Synchronization

Ideal case for parallelism:
- no resources shared between threads
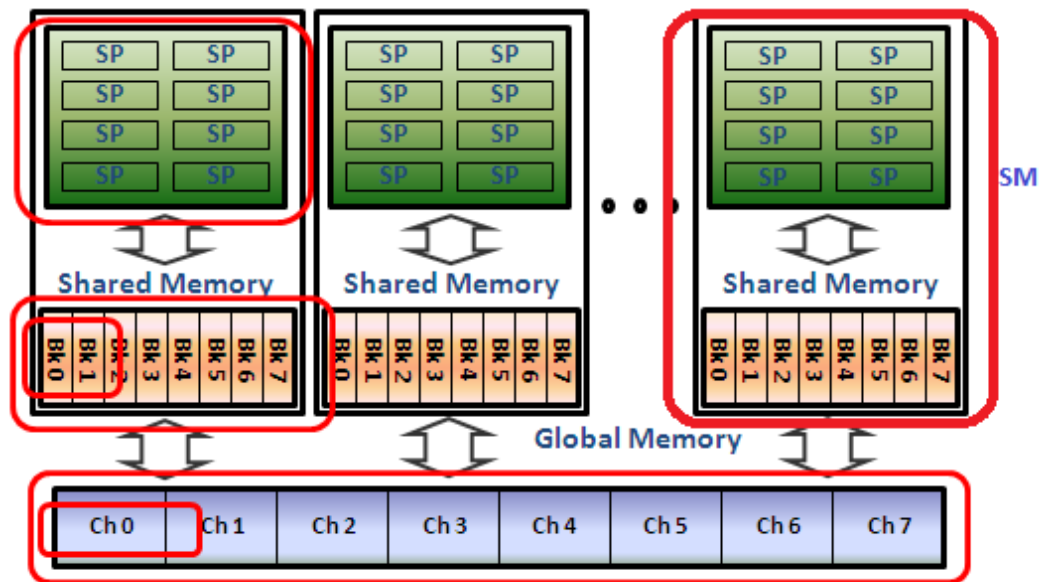- no communication between threads

Many algorithms that require just a little bit of resource sharing can still be accelerated by massive parallelism of GPU

# Examples needing synchronization

(1) Parallel BFS
(2) Summing a list of numbers
(3) Loading data into a GPU's shared memory

# __syncthreads()



- `__syncthreads()` synchronizes all threads in a block.

- Remember that shared memory is per block. Every block that is launched will have to allocate shared memory for its own itself on its resident SM.

- This __synchthreads() call is very useful for kernels using shared memory.

# Atomic instructions: motivation

Two threads try to increment variable x=42  concurrently.
Final value should be 44.
Possible execution order:
```
thread 0 load x (=42) into register r0
thread 1 load x (=42) into register r1
thread 0 increment r0 to 43
thread 1 increment r1 to 43
thread 0 store r0 (=43) into x
thread 1 store r1 (=43) into x
```

Actual final value of x: 43

:(

# Atomic instructions

- An atomic instruction executes as a single unit, cannot be interrupted.

- Serial access

# Atomic instructions on CUDA

```
atomic{Add, Sub, Exch, Min, Max, Inc, Dec, CAS,
        And, Or, Xor}
```

Syntax: `atomicAdd(float *address, float val)`

Work in both global and shared memory!

# (Synchronization) budget advice

Do more cheap things and fewer expensive things!

Example: computing sum of list of numbers

Naive:
each thread atomically increments each number to accumulator in global memory

# Sum example

Smarter solution:
- each thread computes its own sum in register
- use warp shuffle (next slide) to compute sum over warp
- each warp does a single atomic increment to accumulator in global memory
- Reduce number of atomic instructions by a factor of 32 (warp size)

# Warp-synchronous programming

What if I only need to synchronize between all threads in a warp?
Warps are already synchronized!

Can reduce `__syncthreads()` calls

# Warp shuffle

Read value of register from another thread in warp.

```
int __shfl(int var, int srcLane, int width=warpSize)
```

Extremely useful to compute sum of values across a warp.

First available on Kepler (no Fermi, only CC >= 3.0)

# Quick Aside: blur_v from Lab 1

```
blur_device.cu              ×
 8
 9   #include <cuda_runtime.h>
10
11   #include "blur_device.cuh"
12
13
14   __global__
15   void cudaBlurKernel(const float *raw_data, const float *blur_v, float *out_data,
16       int n_frames, int blur_v_size) {
17
18       // TODO: Fill in the implementation for the GPU-accelerated convolution.
19       //
20       // It may be helpful to use the information in the lecture slides, as well
21       // as the CPU implementation, as a reference.
22   }
23
```

```
12       // CPU convolution
13       {
14           for (int i = 0; i < GAUSSIAN_SIZE; i++) {
15               for (int j = 0; j <= i; j++)
16                   output_data_host[i] += input_data[i - j] * blur_v[j];
17           }
18           for (int i = GAUSSIAN_SIZE; i < n_frames; i++) {
19               for (int j = 0; j < GAUSSIAN_SIZE; j++)
20                   output_data_host[i] += input_data[i - j] * blur_v[j];
21           }
22       }
23
```

Shared memory is great place to put blur_v.

1) blur_v is relatively small and easily fits in shared memory.

2) Every thread reads from blur_v

3) Stride 0 access. No bank conflicts – all values broadcasted for i > GAUSSIAN_SIZE (majority of threads)

# Lab 2

(1) Questions on latency hiding, thread divergence, coalesced memory access,  bank conflicts, instruction dependencies

(2) What you actually have to do: Need to comment on all non-coalesced memory accesses and bank conflicts in provided kernel code. Lastly, improve the matrix transpose kernel by using cache and memory optimizations.

# Matrix Transpose



```
C:\Windows\system32\cmd.exe

C:\Users\sunbo\Desktop\lab2>transpose
Size 512 naive CPU: 0.717173 ms
Size 512 GPU memcpy: 0.049180 ms
Size 512 naive GPU: 0.035495 ms
Size 512 shmem GPU: 0.013257 ms
Size 512 optimal GPU: 0.014113 ms

Size 1024 naive CPU: 4.053718 ms
Size 1024 GPU memcpy: 0.068424 ms
Size 1024 naive GPU: 0.013685 ms
Size 1024 shmem GPU: 0.014113 ms
Size 1024 optimal GPU: 0.013685 ms

Size 2048 naive CPU: 42.969670 ms
Size 2048 GPU memcpy: 0.038489 ms
Size 2048 naive GPU: 0.016678 ms
Size 2048 shmem GPU: 0.022666 ms
Size 2048 optimal GPU: 0.014113 ms

Size 4096 naive CPU: 230.006496 ms
Size 4096 GPU memcpy: 0.038489 ms
Size 4096 naive GPU: 0.012402 ms
Size 4096 shmem GPU: 0.022666 ms
Size 4096 optimal GPU: 0.026942 ms
```

An interesting IO problem, because you have a stride 1 access and a stride n access. Not a trivial access pattern like "blur_v" from Lab 1.

Transpose is just a fancy `memcpy`, so `memcpy` provides a great performance target.

Note: This example output is for a clean project without the shmem and optimal kernels completed. Your final output should show a decline in kernel time for the different kernels.

# Matrix Transpose

```
__global__
void naiveTransposeKernel(const float *input, float *output, int n) {
 // launched with (64, 16) block size and (n / 64, n / 64) grid size
 // each block transposes a 64x64 block

 const int i = threadIdx.x + 64 * blockIdx.x;
 int j = 4 * threadIdx.y + 64 * blockIdx.y;
 const int end_j = j + 4;

 for (; j < end_j; j++) {
        output[j + n * i] = input[i + n * j];
 }
}
```

# Shared memory & matrix transpose

Idea to avoid non-coalesced accesses:
● Load from global memory with stride 1
● Store into shared memory with stride x
● `__syncthreads()`
● Load from shared memory with stride y
● Store to global memory with stride 1

Choose values of x and y  perform the transpose.

# Example of an SM's shared memory cache

Showing the first 7/32 banks filled with arbitrary integers

# Example of an SM's shared memory cache



OK: one shared memory access is broadcast in parallel to all the threads in the half-warp.

# Example of an SM's shared memory cache



OK: this access pattern satisfies the main rule as all threads in the half-warp access different shared memory banks.

# Example of an SM's shared memory cache



Not OK: in this case all threads in a half warp access the same bank. The read/writes are bank conflicted, and are performed sequentially. In this example the read/write performance would be 1/8 of maximum.
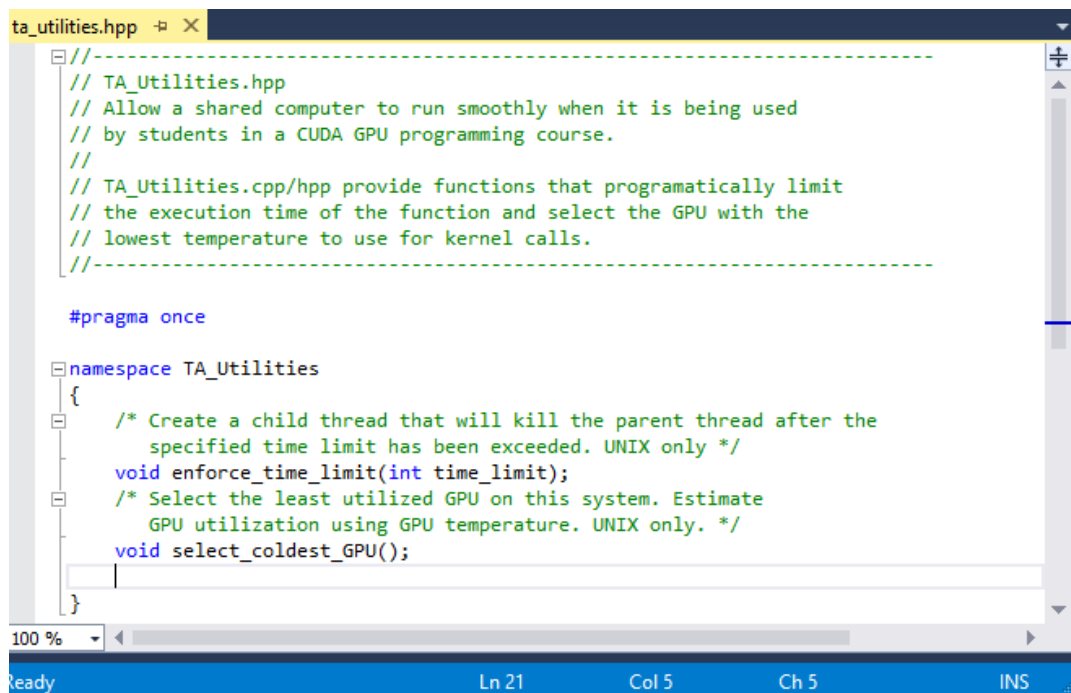
# Avoiding bank conflicts

You can choose x and y to avoid bank conflicts.

A stride n access to shared memory avoids bank conflicts iff `gcd(n, 32) == 1`.

# ta_utils.cpp

- Included in the UNIX version of this set

- Should minimize lag or infinite waits on GPU function calls.

- Please leave these functions in the code if you are using Haru

- Namespace TA_Utilities